

---

# Hands-On: Make Sure You Are Ready

## Step 0: Get what you need.

- Java Development Kit, release 8 or later
- Apache Maven 3.0 or later
- Your favorite IDE or text editor

## Step 1: Download the training project: <https://github.com/GridGain-Demos/imc-essentials-in-90-minutes>.





In-Memory Computing Essentials

November/2020

# Your Instructor: Denis Magda

## → Distributed in-memory system

- ◆ Apache Ignite Committer and PMC Member
- ◆ Head of DevRel at GridGain

## → Java engineering and architecture

- ◆ Java engineering at Oracle
- ◆ Technology evangelism at Sun Microsystems



---

# Training Flow

- Introduction
  - Why In-Memory Computing? (5 mins)
  - Apache Ignite, Brief Overview (15 mins)
- Data Partitioning (25 mins)
- Break (5 mins)
- Affinity Co-location (25 mins)
- Co-located Computations (25 mins)



# Introduction

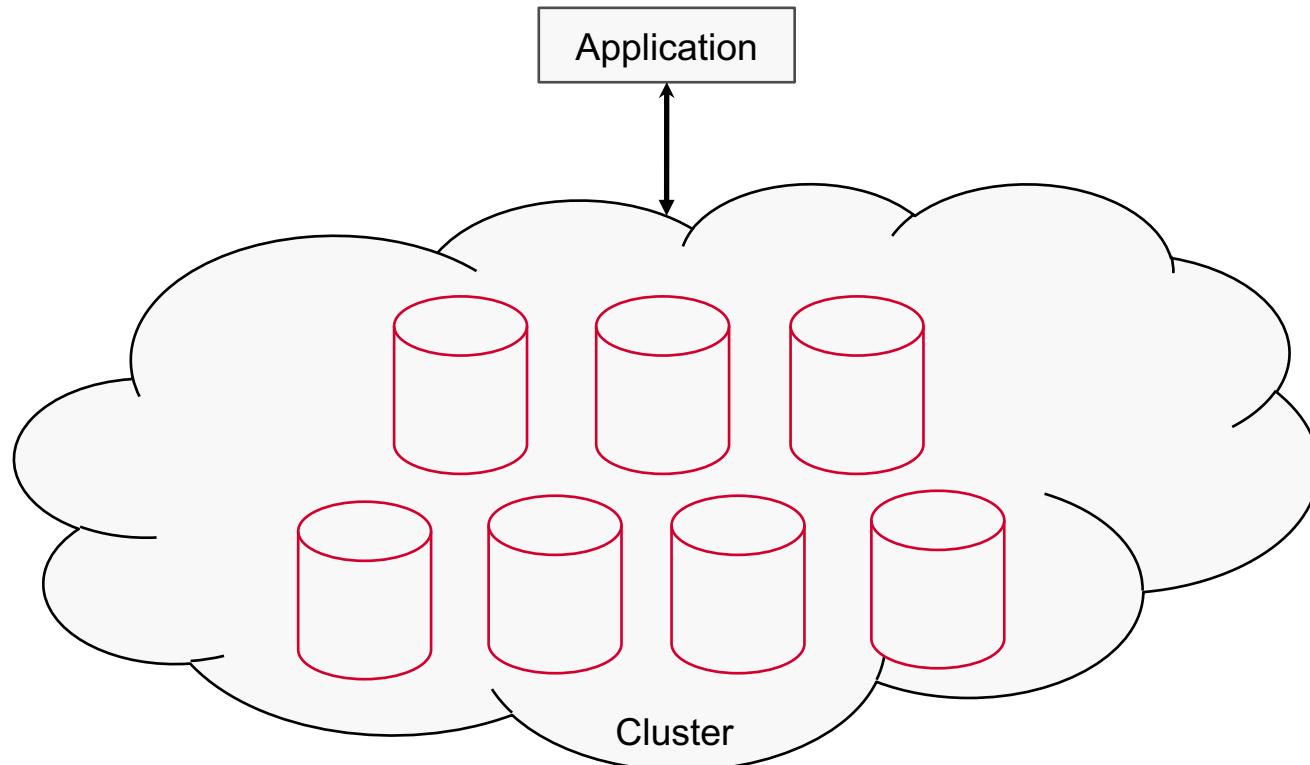
Why In-Memory Computing?

---

# Speed and Scale



# In-Memory Software Is Distributed



# Comparing System-Event Latencies

System Event	Actual Latency	Scaled Latency
One CPU cycle	0.4 ns	1 s
Level 1 cache access	0.9 ns	2 s
Level 2 cache access	2.8 ns	7 s
Level 3 cache access	28 ns	1 min
Main memory access (DDR DIMM)	~100 ns	4 min
Intel Optane DC persistent memory access	~350 ns	15 min
Intel Optane DC SSD I/O	<10 µs	7 hrs
NVMe SSD I/O	~25 µs	17 hrs
SSD I/O	50 – 150 µs	1.5 – 4 days
Rotational disk I/O	1 – 10 ms	1 – 9 months
Internet: SF to NY	65 ms	5 years



# Memory Versus Disk Latency

System Event	Actual Latency	Scaled Latency
One CPU cycle	0.4 ns	1 s
Level 1 cache access	0.9 ns	2 s
Level 2 cache access	2.8 ns	7 s
Level 3 cache access	28 ns	1 min
Main memory access (DDR DIMM)	~100 ns	4 min
Intel Optane DC persistent memory access	~350 ns	15 min
Intel Optane DC SSD I/O	< 10 µs	7 hrs
NVMe SSD I/O	~25 µs	17 hrs
SSD I/O	50 – 150 µs	1.5 – 4 days
Rotational disk I/O	1 – 10 ms	1 – 9 months
Internet: SF to NY	65 ms	5 years



---

# Poll 1

What type of in-memory software do you use?

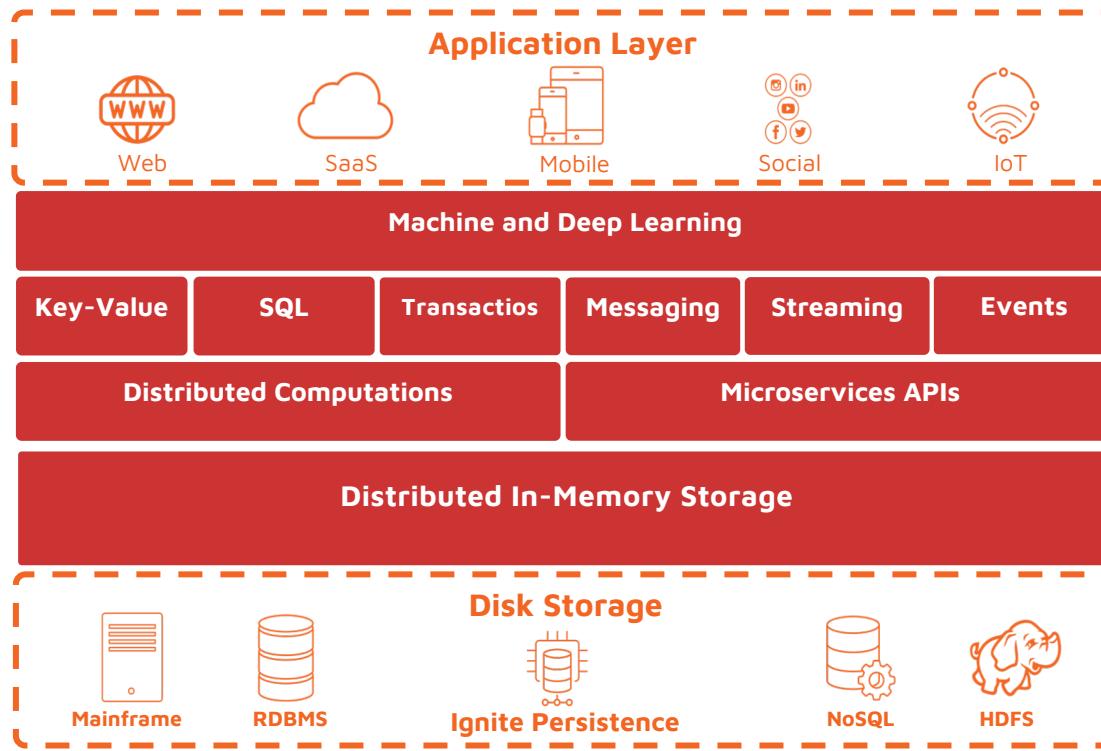
- Caches — Redis, Memcached,...
- Databases — MemSQL, VoltDB,...
- Grids and platforms — Ignite, Hazelcast, Gemfire,...
- Not there yet



# Apache Ignite

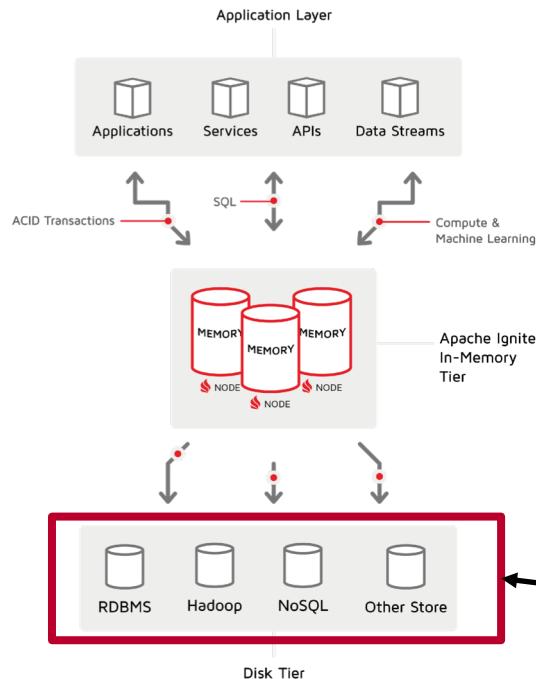
Brief Overview

# Apache Ignite Platform

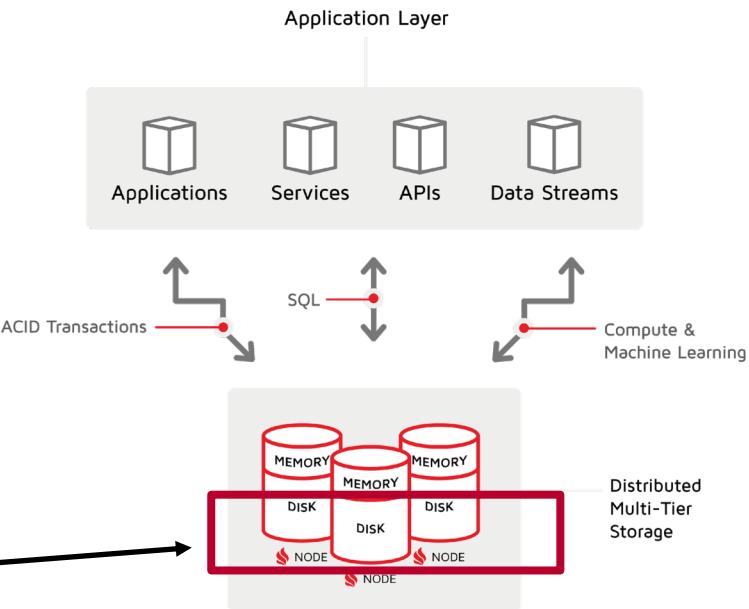


# How Ignite Is Used

## Ignite as a Cache



## Ignite as a Database



difference





## Of more than 350 Apache Projects

### Top 10 Dev Mailing Lists

1.  Flink

...

8.  apache Ignite™

9. APACHE  ARROW

10.  Apache Airflow

### Top 5 User Mailing Lists

1.  Flink

2.  lucene

3.  apache Ignite™

4.  Apache Tomcat

5.  kafka



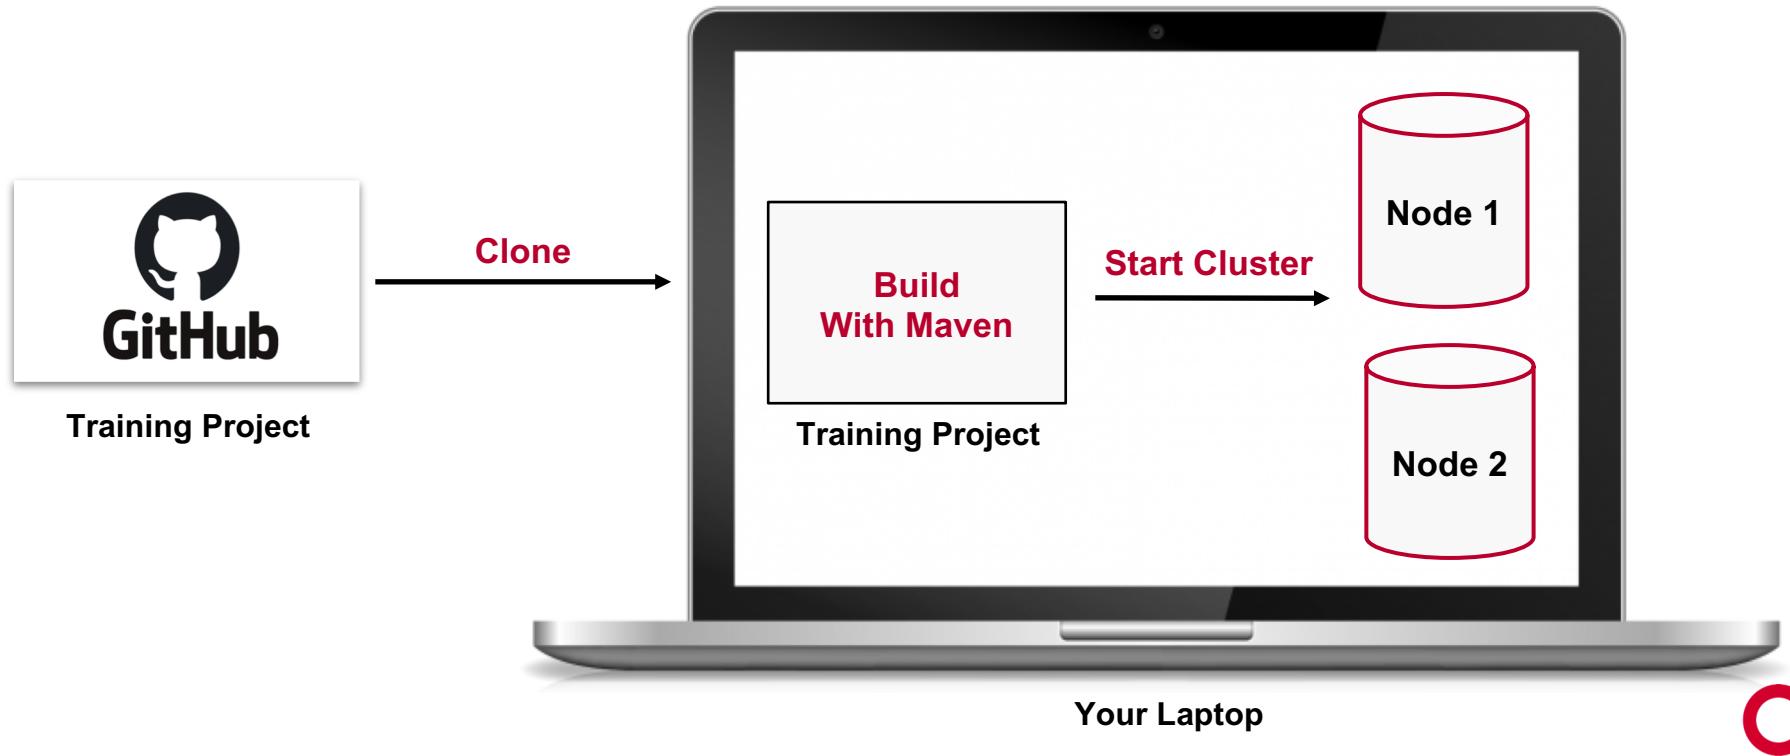
*Apache Ignite is used by thousands of companies.*



Bloomberg



# Hands-On: What You'll Do Next



# Hands-On: Start the Cluster

**Step 0:** Get what you need.

- Java Development Kit, version 8 or later
- Apache Maven 3.0 or later
- Your favorite IDE or text editor

**Step 1:** Download the training project:

<https://github.com/GridGain-Demos/imc-essentials-in-90-minutes>.

**Step 2:** Unzip the project and go to its root directory.



# Hands-On: Start the Cluster

Step 3: Build the core module of the project.



```
mvn clean package -P core
```

Step 4: Start the first cluster node.



```
java -cp libs/core.jar training.ServerStartup
```

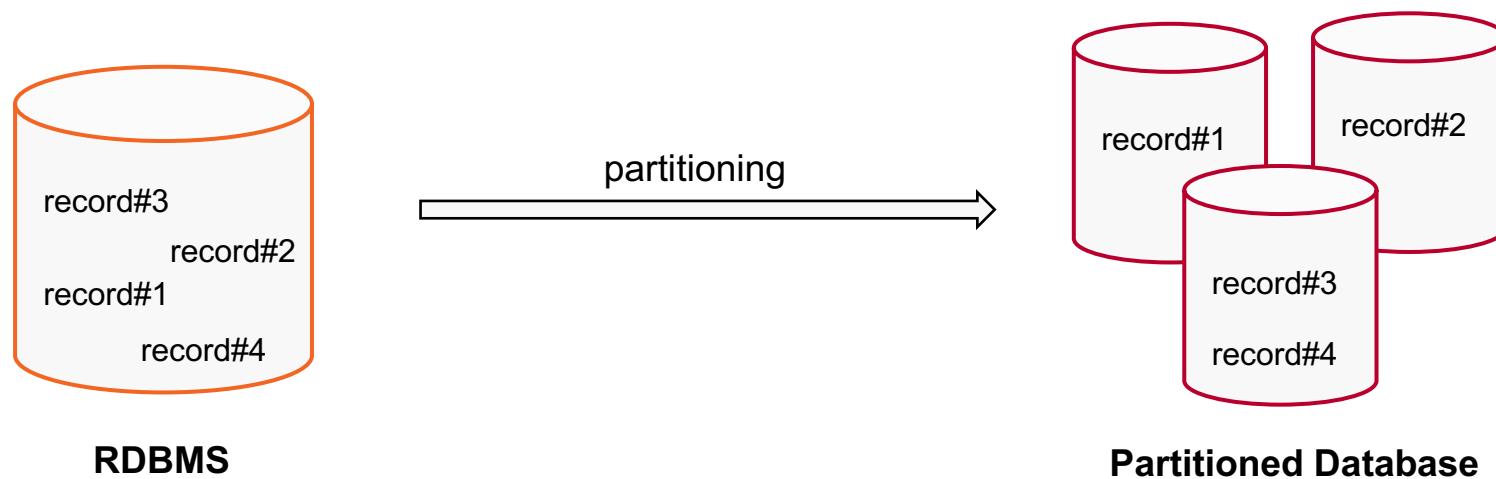
Step 5: In another terminal window, start the second node.



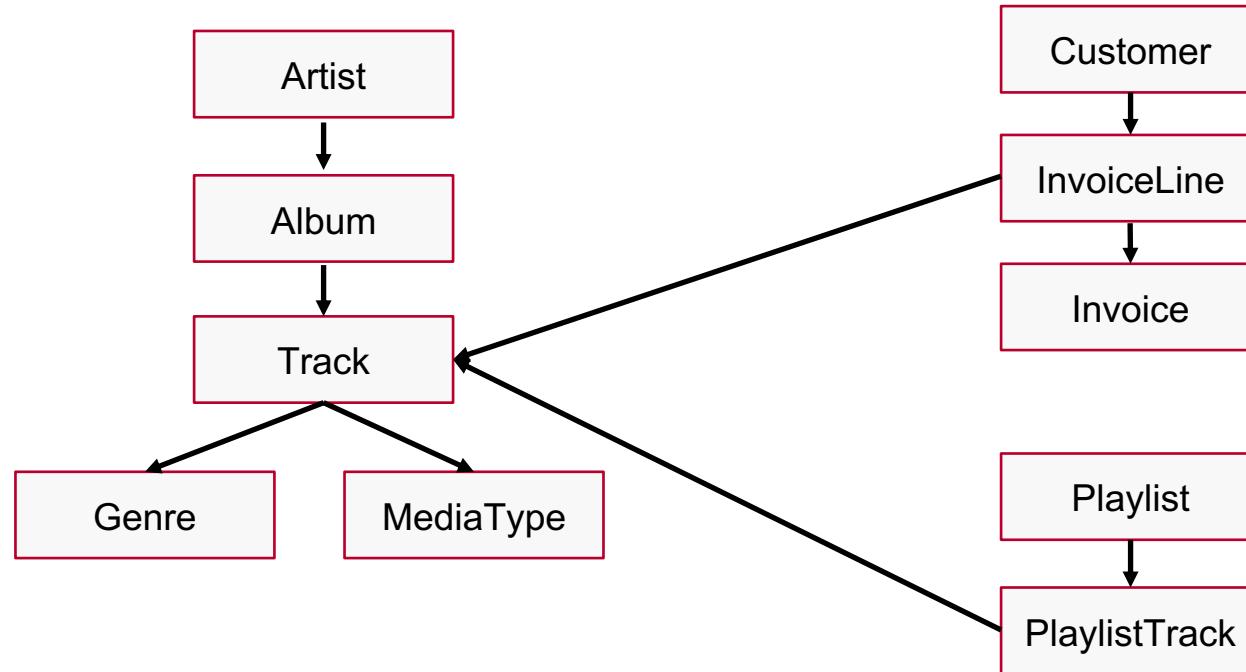
# Data Partitioning

Using all of the cluster's memory and CPUs

# Partitioning and Horizontal Scalability



# Chinook Database: Digital Media Store

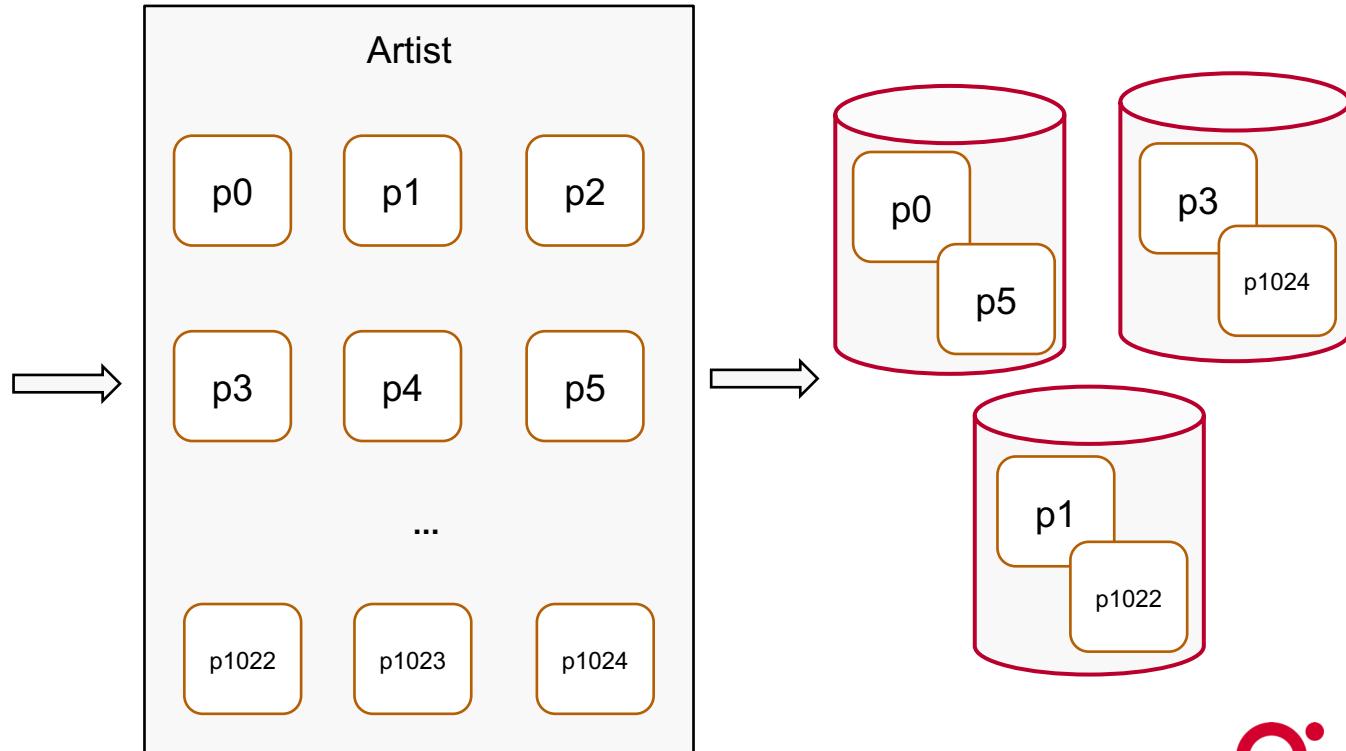


<https://github.com/lerocha/chinook-database>



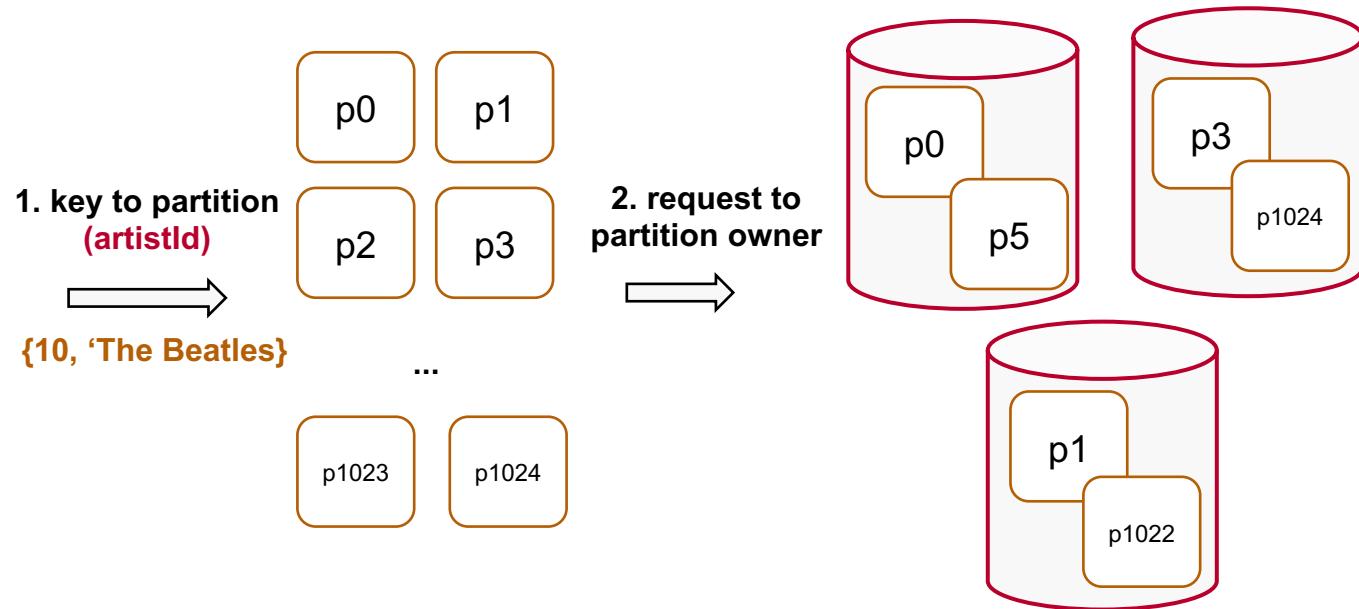
# Partitions

```
CREATE TABLE Artist {  
...  
}
```



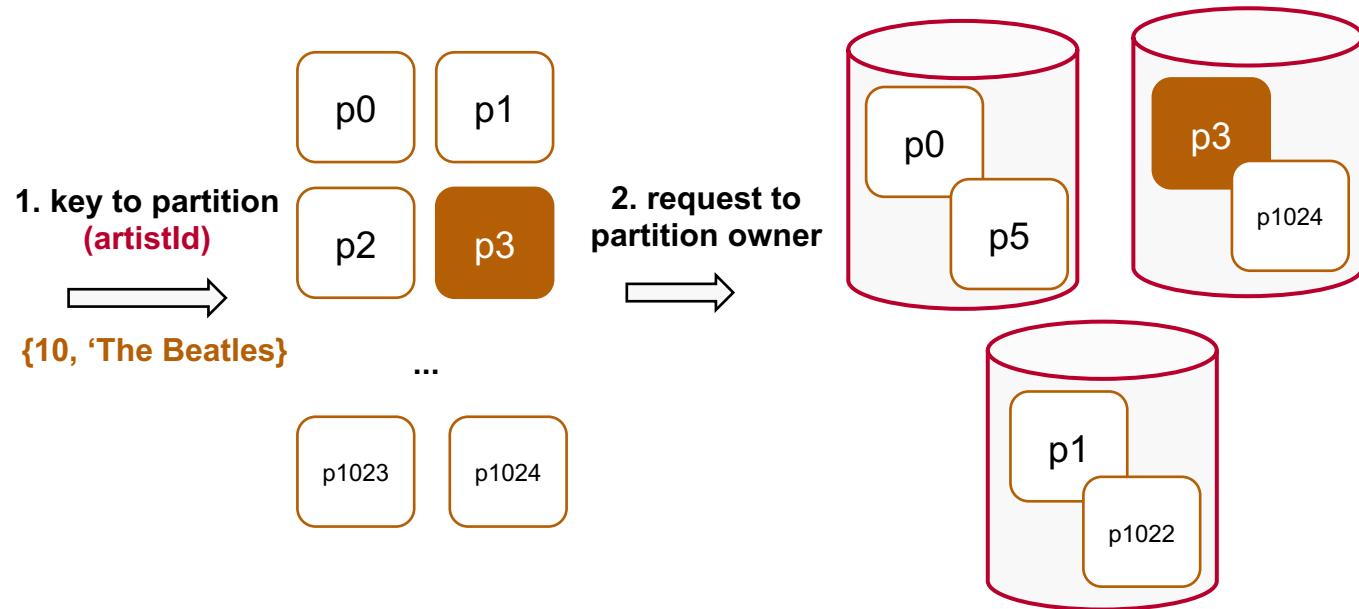
# Record->Partition->Node Mapping

```
INSERT INTO Artist  
(ArtistId, Name)  
VALUES (10, 'The Beatles')
```

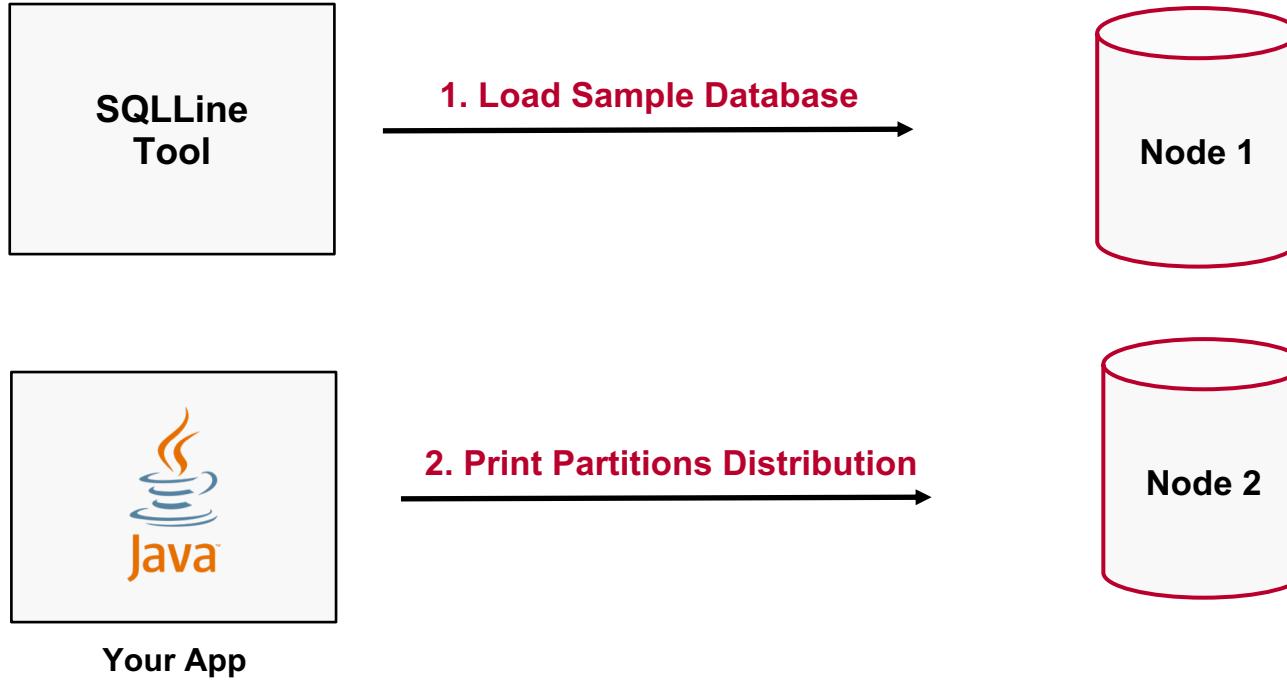


# Record->Partition->Node Mapping

```
INSERT INTO Artist  
(ArtistId, Name)  
VALUES (10, 'The Beatles')
```



# Hands-On: What You'll Do Next



# Hands-On: Load the Chinook Database

Step 1: Start SQLLine tool.

```
java -cp libs/core.jar sqlline.SqlLine
```

Step 2: Use SQLLine to connect to the cluster.

```
!connect jdbc:ignite:thin://127.0.0.1/ ignite ignite
```



---

# Hands-On: Load the Chinook Database

Step 3: Load the database.

```
!run config/media_store.sql
```

Step 4: List the created tables.

```
!tables
```



# Hands-On: Print Partitions' Distribution

Step 1: Update the oreilly.training.KeyValueApp logic.

```
● ● ●

Affinity<Integer> affinity = ignite.affinity("Artist");

int part = affinity.partition(artistKey);
UUID nodeId = affinity.mapPartitionToNode(part).id();

System.out.println(artist + " [part=" + part + ", node=" + nodeId);
```



# Hands-On: Print Partitions Distribution

Step 2: Build a version of the project that has application classes:



```
mvn clean package -P apps
```

Step 3: Start the application and check partitions' distribution



```
java -cp libs/apps.jar training.KeyValueApp
```

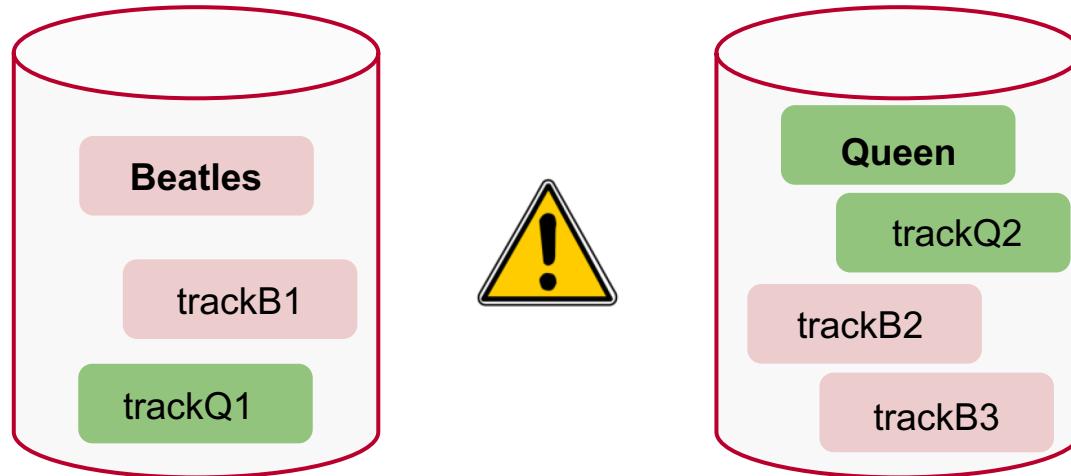


# Five-Minute Break

# Affinity Co-location

Reducing network utilization for complex requests

# Random Data Distribution



Queen-related data



Beatles-related data



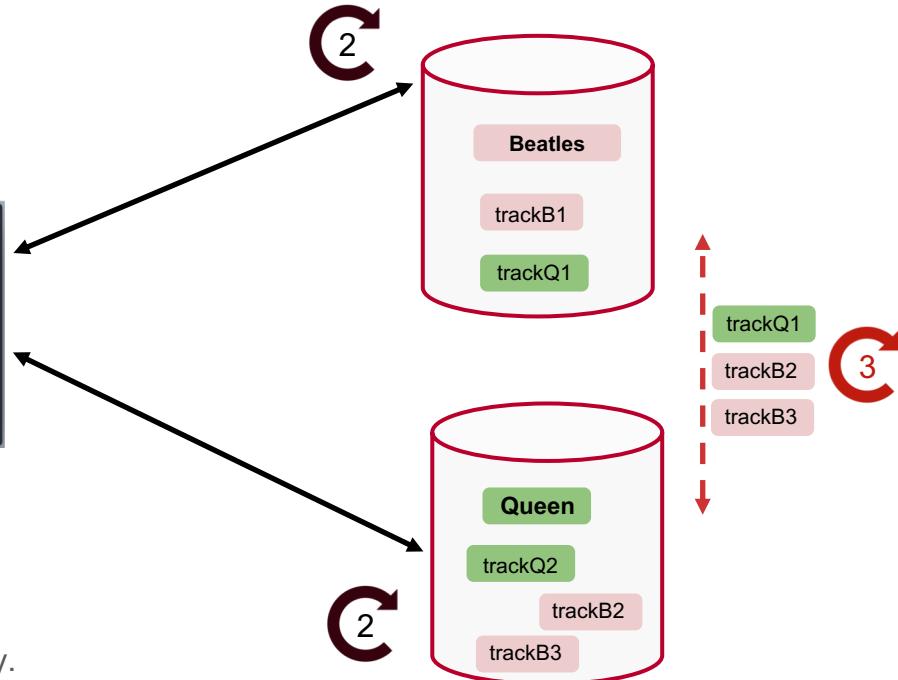
# Data Is Shuffled During the JOIN phase

Application



1 & 4

```
...  
SELECT track.trackId, track.name as track_name, genre.name as genre,  
artist.name as artist, MAX(milliseconds / (1000 * 60)) as duration  
FROM track LEFT JOIN artist ON track.artistId = artist.artistId  
JOIN genre ON track.genreId = genre.genreId  
WHERE track.genreId < 17  
GROUP BY track.trackId, track.name, genre.name, artist.name  
ORDER BY duration DESC LIMIT 20;
```



1. The application starts executing the query.
2. The query is executed on the cluster nodes.
3. **Data is shuffled between the nodes on JOIN.**
4. The result set is reduced.

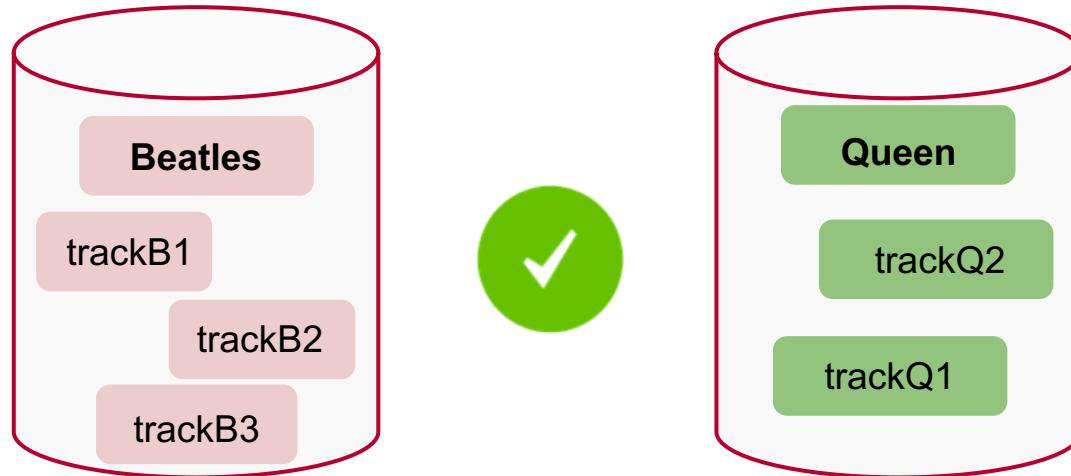


# Disk Versus Network Latency

System Event	Actual Latency	Scaled Latency
One CPU cycle	0.4 ns	1 s
Level 1 cache access	0.9 ns	2 s
Level 2 cache access	2.8 ns	7 s
Level 3 cache access	28 ns	1 min
Main memory access (DDR DIMM)	~100 ns	4 min
Intel Optane DC persistent memory access	~350 ns	15 min
Intel Optane DC SSD I/O	<10 µs	7 hrs
NVMe SSD I/O	~25 µs	17 hrs
SSD I/O	50-150 µs	1.5 – 4 days
Rotational disk I/O	1 – 10 ms	1 – 9 months
Internet: SF to NY	65 ms	5 years



# Co-located Data Distribution



Queen-related data



Beatles-related data



# How to Group Related Data

```
CREATE TABLE Artist
(
    ArtistId INT,
    Name VARCHAR(120),
    PRIMARY KEY (ArtistId)
);
```

```
CREATE TABLE Track
(
    TrackId INT,
    Name VARCHAR(200),
    AlbumId INT,
    ArtistId INT,
    GenreId INT,
    Composer VARCHAR(220),
    Milliseconds INT,
    PRIMARY KEY (TrackId, ArtistId)
) WITH affinityKey=ArtistId;
```



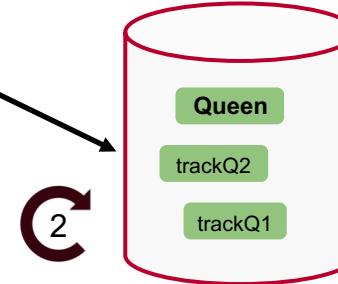
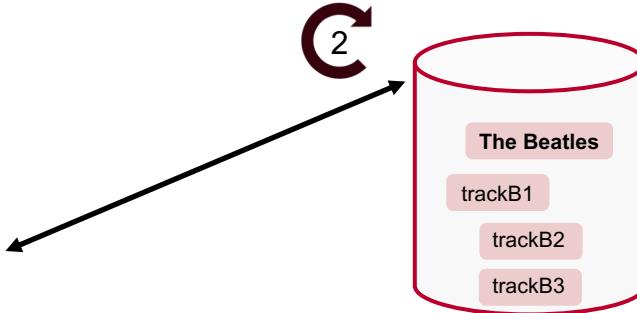
# JOINS with Co-Located Data

Application



1 & 4

```
...  
SELECT track.trackId, track.name as track_name, genre.name as genre,  
artist.name as artist, MAX(milliseconds / (1000 * 60)) as duration  
FROM track LEFT JOIN artist ON track.artistId = artist.artistId  
JOIN genre ON track.genreId = genre.genreId  
WHERE track.genreId < 17  
GROUP BY track.trackId, track.name, genre.name, artist.name  
ORDER BY duration DESC LIMIT 20;
```



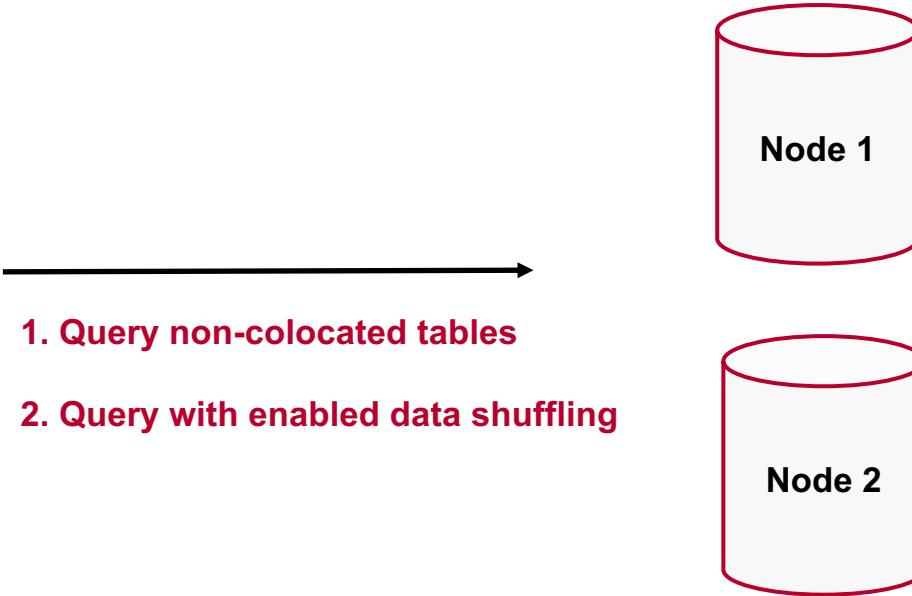
1. The application starts executing the query.
2. The query is executed on the cluster nodes.
3. The result set is reduced.



# Hands-On: What You'll Do Next



1. Query non-colocated tables  
2. Query with enabled data shuffling



# Hands-On: JOIN Without Data Co-location

Step 1: Make sure that you are connected to the cluster via SQLline.

```
● ● ●  
java -cp libs/core.jar sqlline.SqlLine  
  
!connect jdbc:ignite:thin://127.0.0.1/ ignite ignite
```

Step 2: Get the top 20 longest tracks.

```
● ● ●  
!run sql/top_20_longest_tracks.sql
```



# Hands-On: JOINs Without Data Co-location

**Step 3:** Get the top 20 with details about the artists and identify data inconsistencies:



```
!run sql/top_20_longest_tracks_with_authors.sql
```

**Step 4:** Allow data shuffling for JOINs and run the query again:



```
!close
```

```
!connect jdbc:ignite:thin://127.0.0.1?distributedJoins=true ignite ignite
```



---

# Demo by Instructor: Co-locating Related Data

**Step 1:** Co-locate tracks with artists.

- Change Tracks' primary key to **PRIMARY KEY (TrackId, ArtistId)**.
- Set Track's affinity key to **affinityKey=ArtistId**.

**Step 2:** Clean the Ignite work directory and restart the cluster.

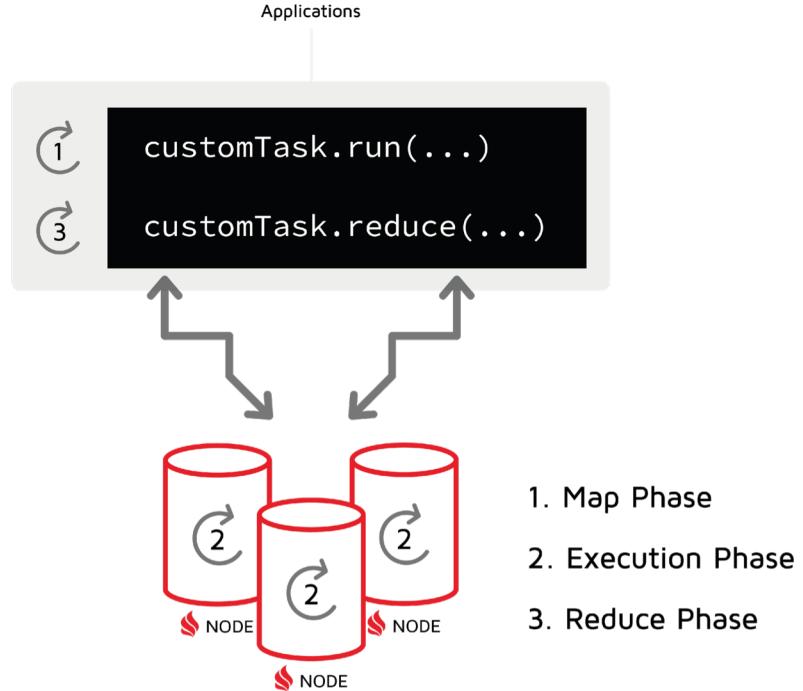
**Step 3:** Reconnect with SQLLine and run the last SQL query.



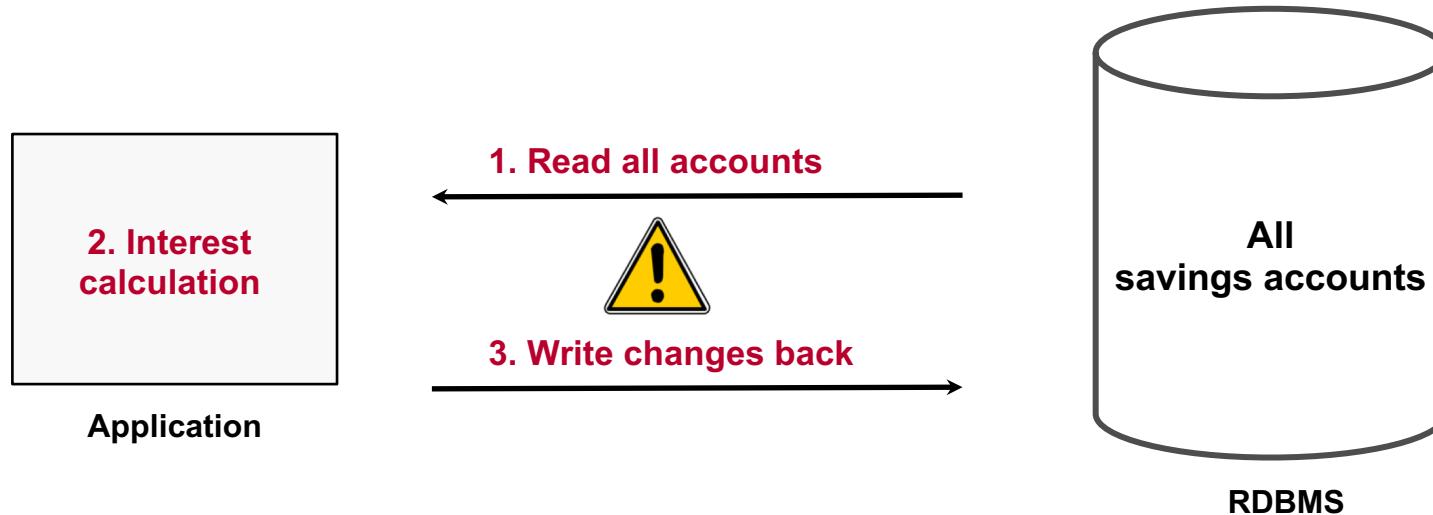
# Co-located Computations

Executing data-intensive logic on cluster nodes

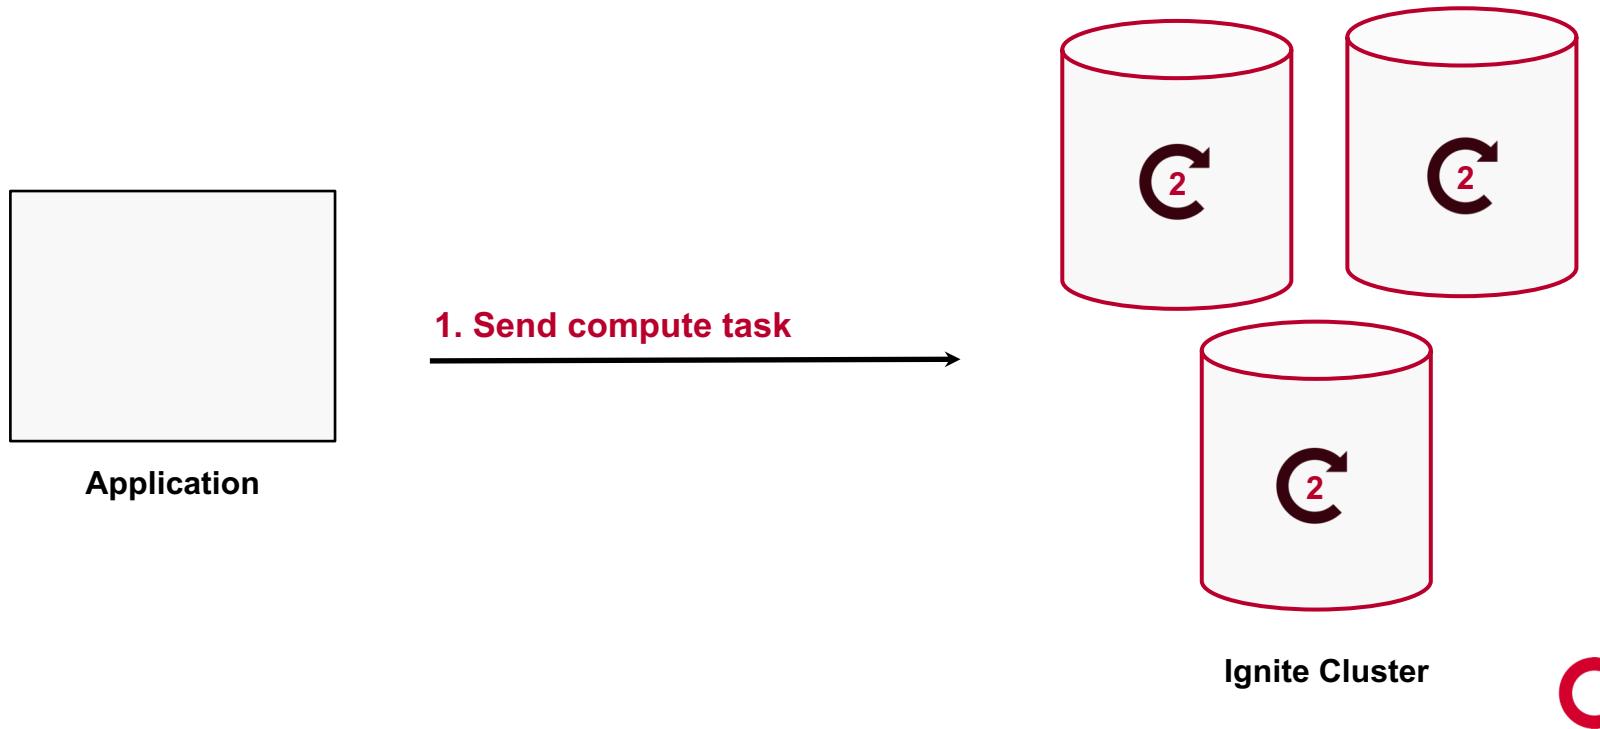
# Executing Custom Logic in the Cluster



# Case With Millions Savings Accounts



# Case With Millions Savings Accounts



---

# Poll 2

What was the new execution time of the business operation once it was implemented as a compute task?

- 40 - 50 minutes
- 10 - 15 minutes
- 1 - 5 minutes
- Less than a minute

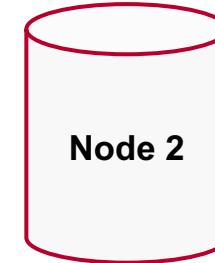
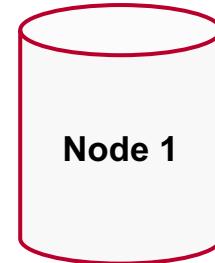


# Hands-On: What You'll Do Next

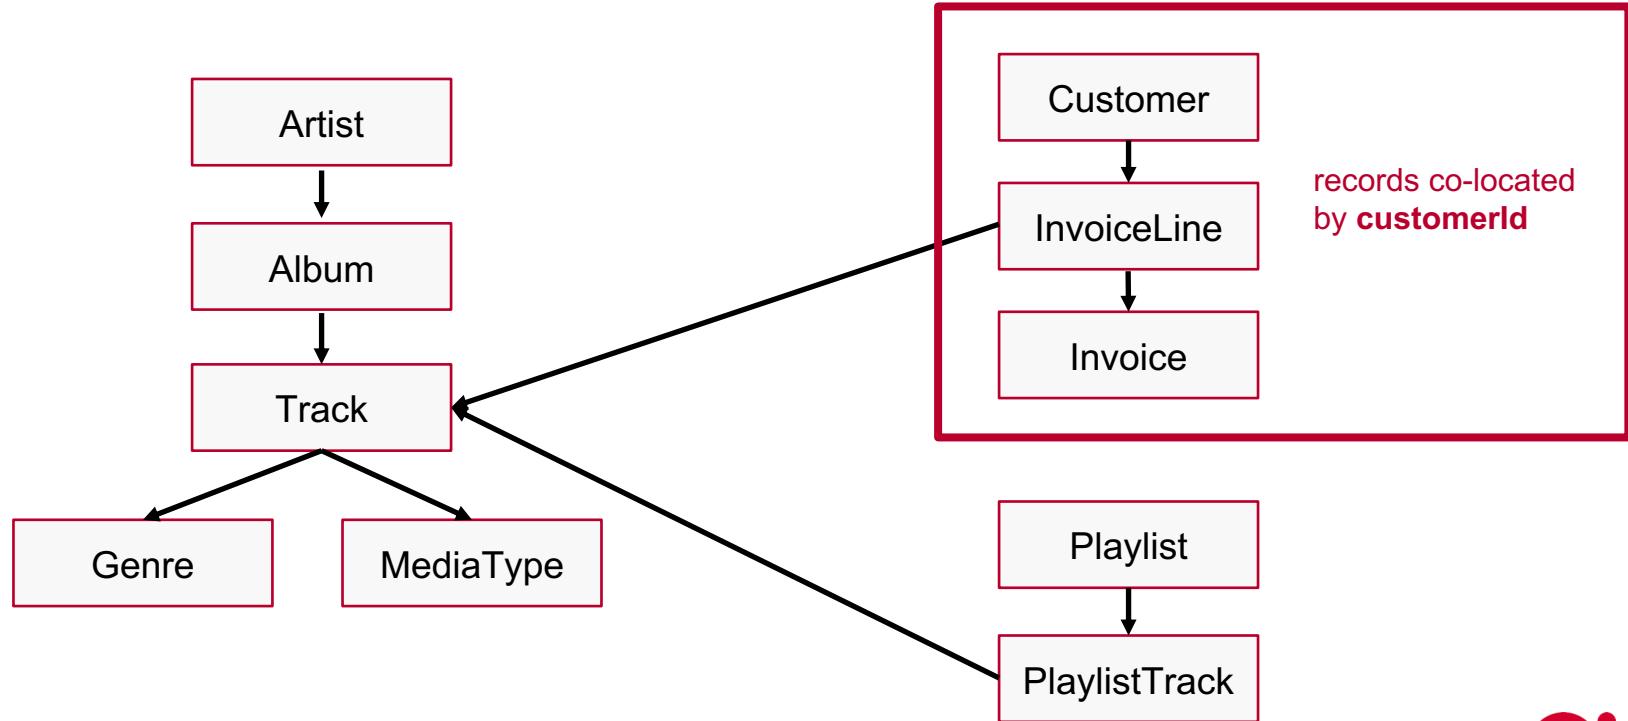


Your App

1. Find top-5 paying listeners
2. Find top-10 paying listeners



# Top-5 Paying Listeners of the Media Store



# Hands-On: Find Top Paying Listeners

Step 1: Execute the Java logic that returns top-5 paying listeners.



```
java -cp libs/apps.jar training.ComputeApp
```

Step 2: Update the source code to return top-10 paying listeners.



# Hands-On: Find the Top-5 Paying Listeners

Step 3: Rebuild the apps' artifact of the project:

```
mvn clean package -P apps
```

Step 4: Execute the application.

The servers will reload the updated version of the computation.

```
java -cp libs/apps.jar training.ComputeApp
```



# Summary

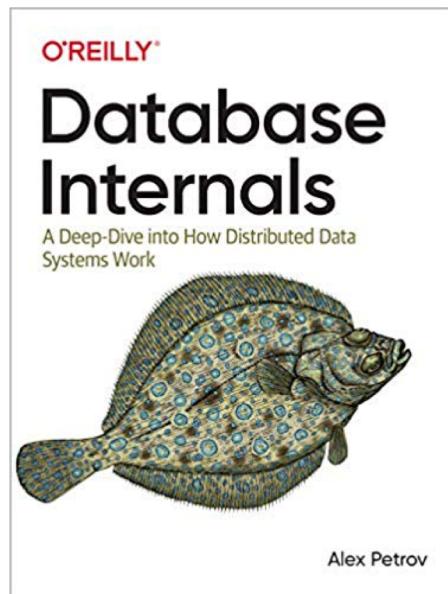
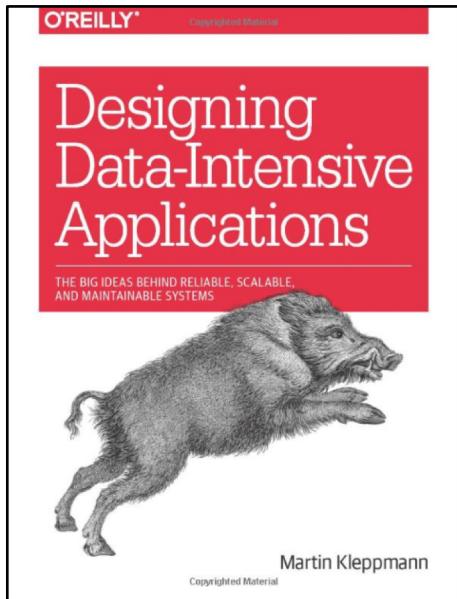
In-memory computing essentials

# It's Not Enough to Put Data in RAM

- **Scale out** to spread data evenly and to balance the load.
  - Understand partitioning basics
  - Know how to tweak default partitioning behavior
- **Avoid** any negative impact that the network might have on performance
  - Co-locate the related data.
  - Run data-intensive logic on the cluster nodes.



# Want to Learn More?



A screenshot of the Apache Ignite website. The header includes the Ignite logo and navigation links for Features, Use Cases, Resources, Community, Events, and Download. The main content area is titled 'Apache Ignite In-Memory Computing Platform' and features a 'Getting Started' button. Below this is a GitHub repository card with statistics: 3,392 stars, 5,771 forks, 280 watches, and 7,851 followers. A descriptive paragraph explains Ignite's capabilities: 'Apache Ignite® is a horizontally scalable, fault-tolerant distributed in-memory computing platform for building real-time applications that can process terabytes of data with in-memory speed.' There are three circular icons representing 'In-Memory Cache', 'In-Memory Data Grid', and 'In-Memory Database', each with a brief description.

 @denismagda





Thank you

